

Colloquium

# Lab: Concurrent and Distributed Systems

## Himeno Benchmark and Mandelbrot Set

Fritz Louis Wilke  
fritz.wilke@tu-dresden.de  
4536116  
M. Sc. Computer Science

# Agenda

## The **Mandelbrot Set**

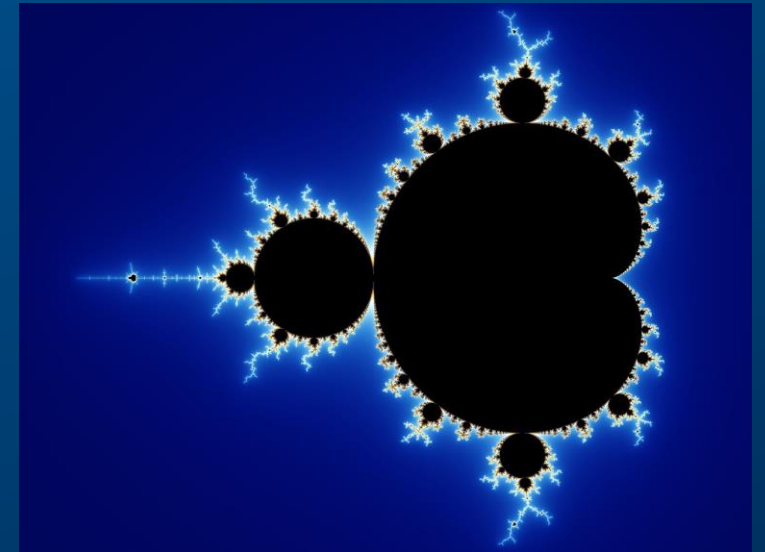
- **Problem**
- (Unoptimized) **Algorithm**
- **Optimizations**
- **Evaluation**

## The **Himeno Benchmark**

- **Problem**
- (Unoptimized) **Algorithm**
- **Optimizations**
- **Evaluation**

## Summary

# The Mandelbrot Set



[https://commons.wikimedia.org/wiki/File:Mandel\\_zoom\\_00\\_mandelbrot\\_set.jpg](https://commons.wikimedia.org/wiki/File:Mandel_zoom_00_mandelbrot_set.jpg)



# Unoptimized Algorithm

## The Mandelbrot Set

```
input: uint rows, uint cols, uint nn
output: char[]
begin
    let img = char[rows][cols]
    for r in 0..rows:
        for c in 0..cols:
            let z = complex(0, 0)
            let n = 0

            while |z| < 2.0 and n < nn:
                z = z*z + complex( c*2 / cols - 1.5, r*2 / rows - 1.0 )
                n = n + 1
            img[r][c] = (n == nn) ? '#' : '.'

    for r in 0..rows:
        for c in 0..cols: print( img[r][c] )
        print( '\n' )
end
```

# Optimizations

## The Mandelbrot Set

### General:

- **Simplifying** calculations
- Working directly with **floats** instead of `std::complex<float>`
- Storing and **reusing** interim results
- Working on **continuous** memory (as image)

### Parallelization:

- **Parallelize** the pixel calculation (inner for loop)
- Provide and collect work **dynamically** via **queues**
- Aligning data structures to fit **cache lines**
- **Each** worker thread has one **input-** and **output queue**

# Parallel Workflow

## The Mandelbrot Set



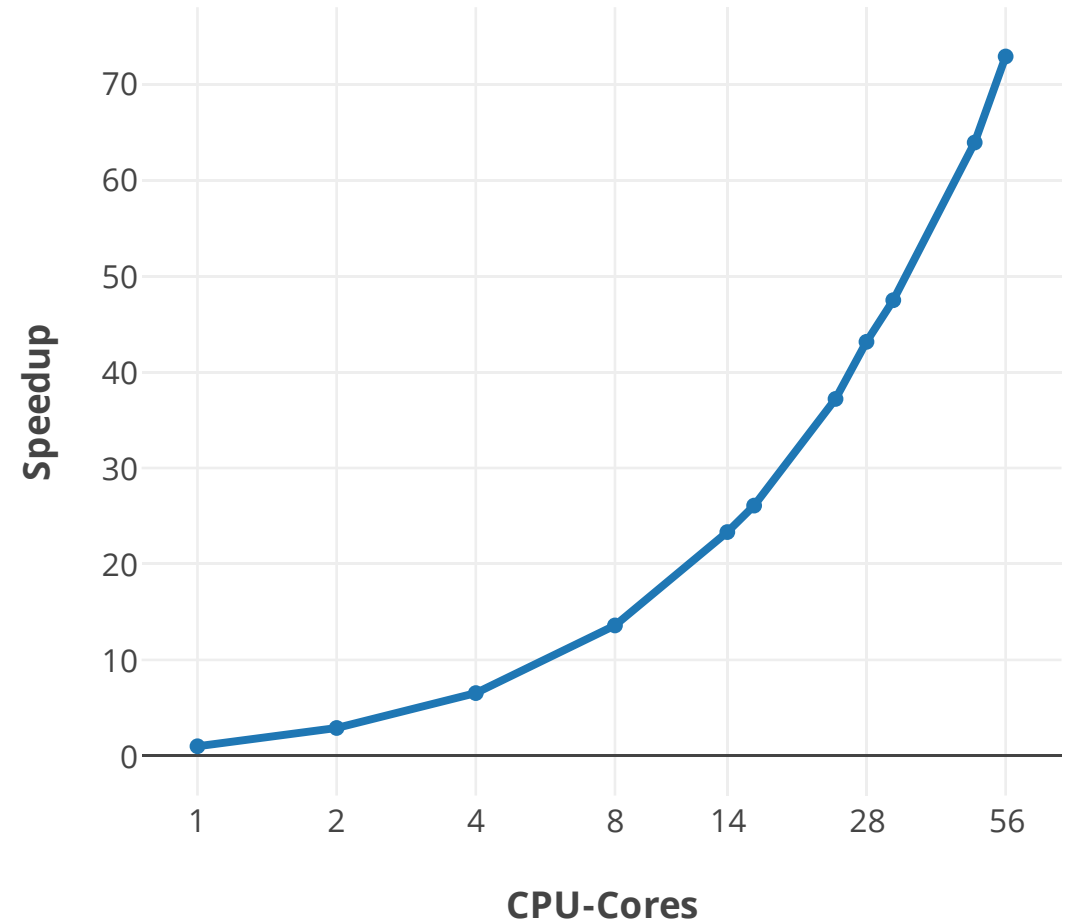
# Evaluation

## The Mandelbrot Set

- Unoptimized code as **reference**
- **C++** and **C#**
- Scaled very well
- No decline after **14** or **28** cores due to local queues
- Parallel code amount over 100% (not realistic)

Lang.	V	Optimization	Time	Abs.	Rel.
C++	0	Unmodified	207050	1x	1x
	1	I/O Threads	2182	86x	83x
	2	+ Simplified	<b>263</b>	<b>716x</b>	73x
C#	0	Unoptimized	63332	1x	1x
	2	Both	452	146x	<b>1798x</b>

Scaling of C++ Version 2





# The Himeno Benchmark

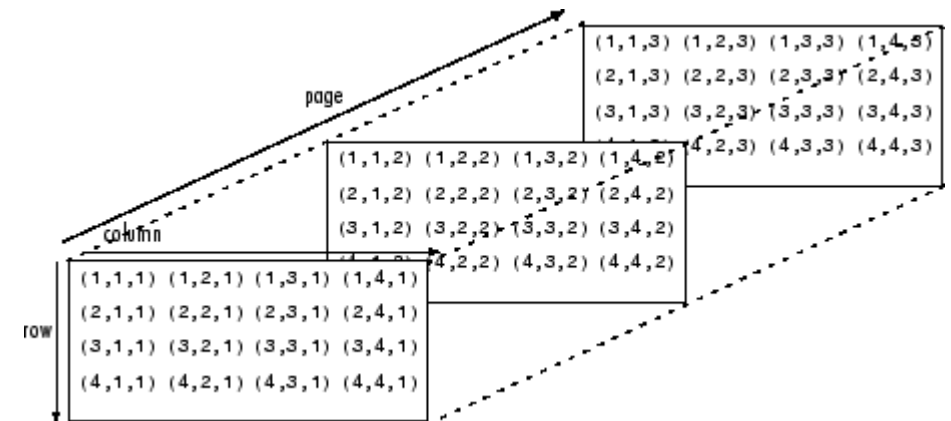


<http://www.isum2017.udg.mx/es/contenido/dr-ryutaro-himeno>

# Problem

## The Himeno Benchmark

- A **benchmarking** program
- Extremely **memory intensive**
- Uses the **Jacobi method** for calculations
- Works on several 4D-Matrices, where first 3 dimensions are same size



<http://www.isum2017.udg.mx/es/contenido/dr-ryutaro-himeno>

# Unoptimized Algorithm

## The Himeno Benchmark

```
input: uint rows, uint cols, uint deps, uint nn
output: float
begin
    initialize_matrices()
    let gosa

    for n in 0..nn:
        gosa = 0.0
        for r in 1..rows-1: // iterate over every voxel
            for c in 1..cols-1:
                for d in 1..deps-1:
                    let ss = jacobi_iteration(r, c, d)
                    gosa += ss*ss
                    wrk2[0][r][c][d] = p[0][r][c][d] + 0.8*ss
                wrk2.copy_to(p) // copy content of p to wrk2

    print( gosa ) // print result (with a precision of 6)
end
```

# Optimizations

## The Himeno Benchmark

### General:

- Drastically **reducing memory usage**:
  - **Substitute** 5 of 7 matrices by constants
  - From 33 to 9 **matrix reads/writes (by 73%)**
  - **Pointer swapping** instead of matrix copying (reduces  $(rows - 2) * (cols - 2) * (deps - 2)$  to 1) per it.
  - **Reduce** matrix sizes by 2 each dimension
- Only calculate the gosa in **last iteration**

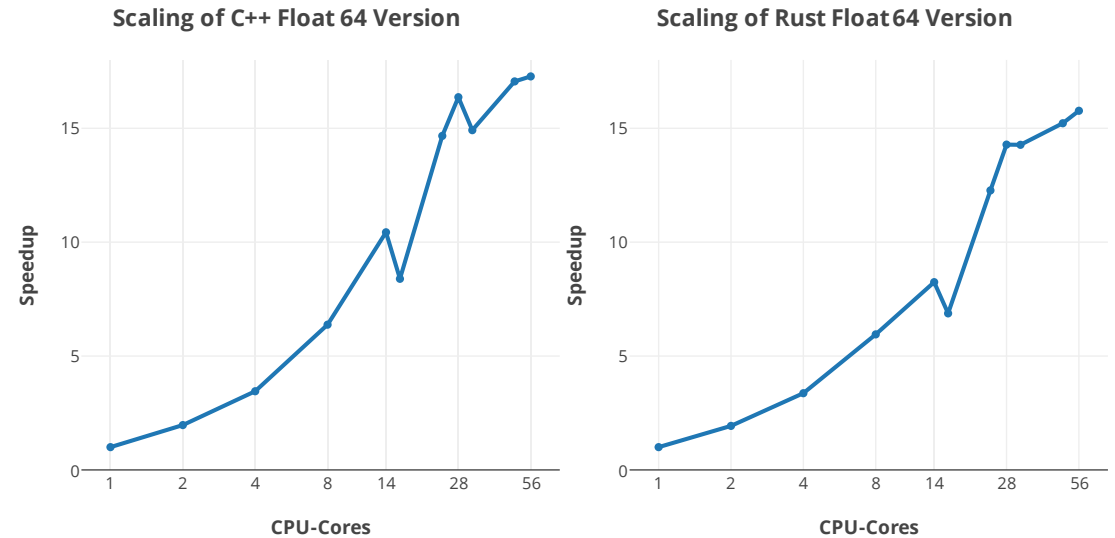
### Parallelization:

- Parallelize on the **rows loop** (outermost, parallelizable loop)
- Threads get work from a **shared atomic variable** containing the next row
- Calculate gosa **sub-sums** in parallel (only F64)

# Evaluation

## The Himeno Benchmark

- Unoptimized code as **reference**
- **C++** and **Rust**
- Scaled ok (w.r.t. the high memory usage)
- Strong decline at **16** and **32** cores
- Parallel code up to **97.37%**

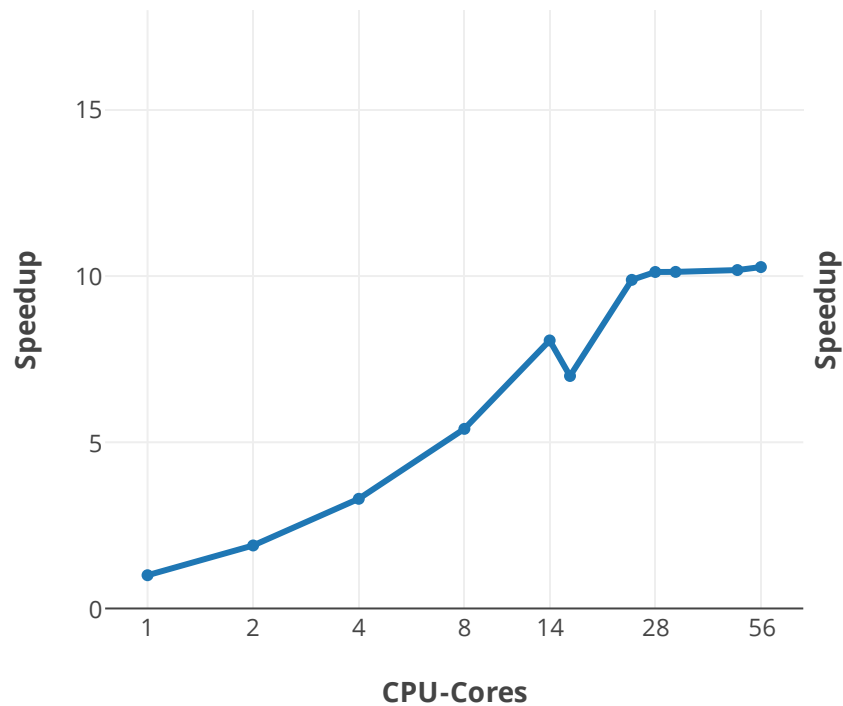


Lang.	FXX	Optimization	Time	Abs.	Rel.	p @ 28	p @ 56
C++	32	Unmodified	171978	1x	1x	0%	0%
		Parallel	34346	5x	10x	93.46%	91.91%
	64	Unmodified	207050	1x	1x	0%	0%
		Parallel	19294	10x	<b>17x</b>	<b>97.37%</b>	<b>95.92%</b>
Rust	64	Unoptimized	289943	1x	1x	0%	0%
		Parallel	<b>7108</b>	<b>43x</b>	16x	96.44%	95.36%

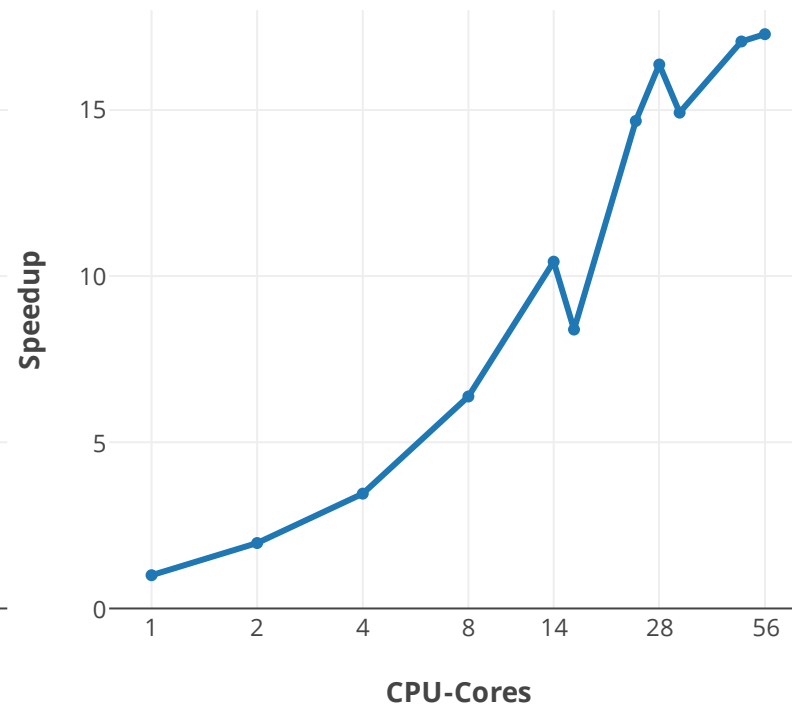
# Evaluation

## The Himeno Benchmark

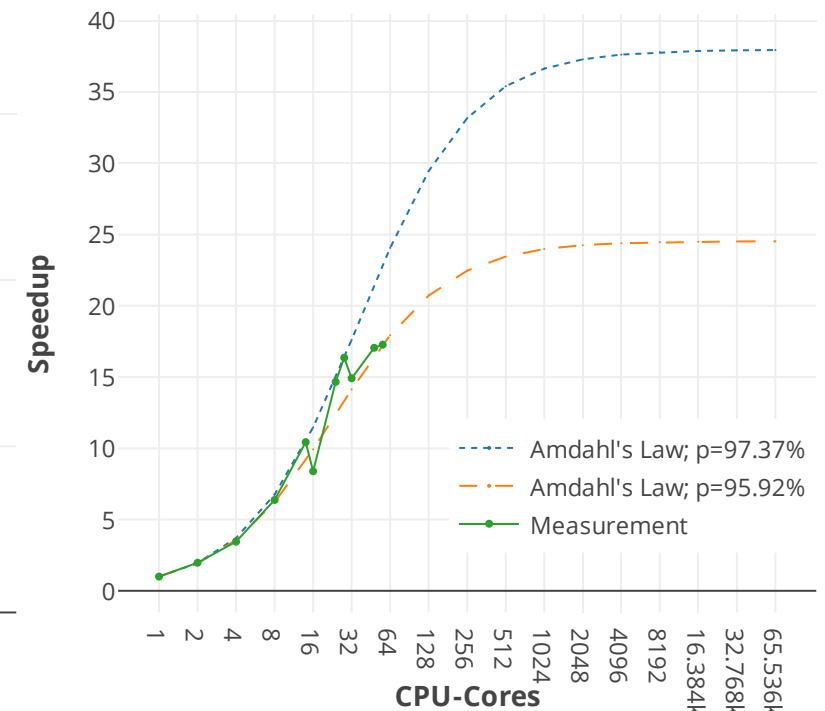
Scaling of C++ Float 32 Version



Scaling of C++ Float 64 Version



Scaling of C++ Float 64 Version



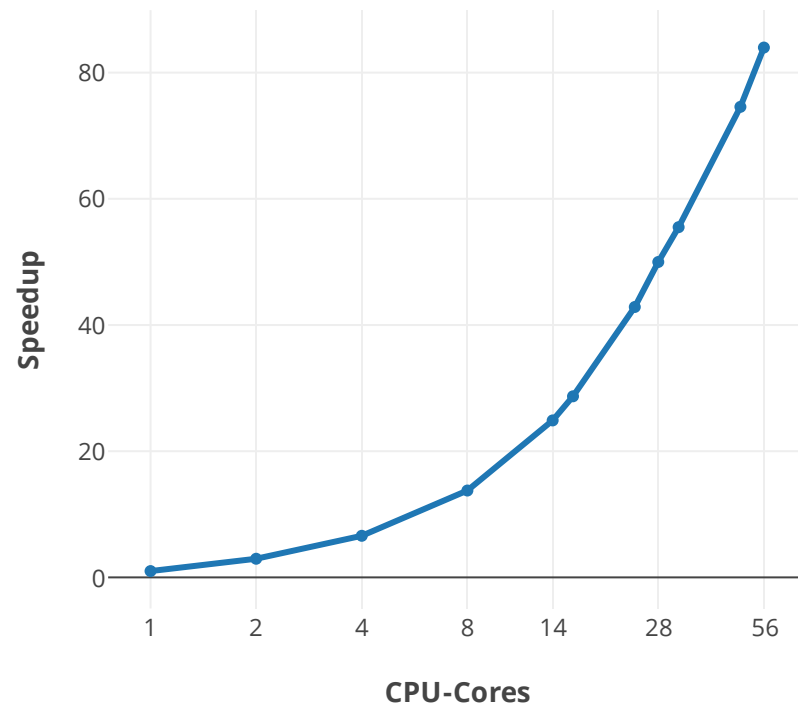
# Summary

- Solved workload distribution with **queues** and **atomic variables**
- **Optimized** and **simplified** the algorithms
- Speedup of **1798** (716 abs.) for the **Mandelbrot Set**, optimal goal was **51**
- Speedup **of 17** (43 abs.) for the **Himeno Benchmark**, optimal goal was **13**
- Learned about **memory throughput** and **cache invalidation** as limiting factor

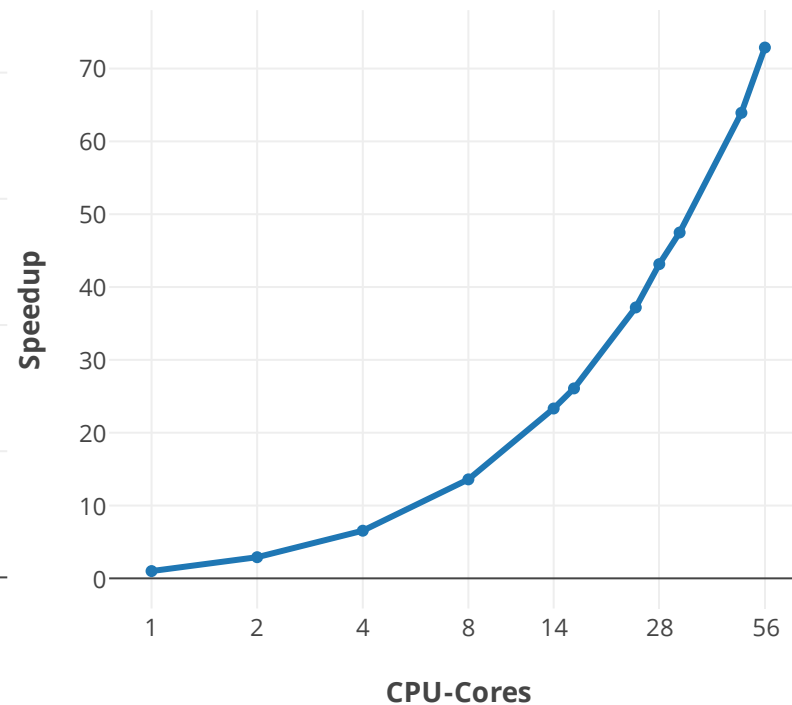
# Scaling of all Versions

## The Mandelbrot Set

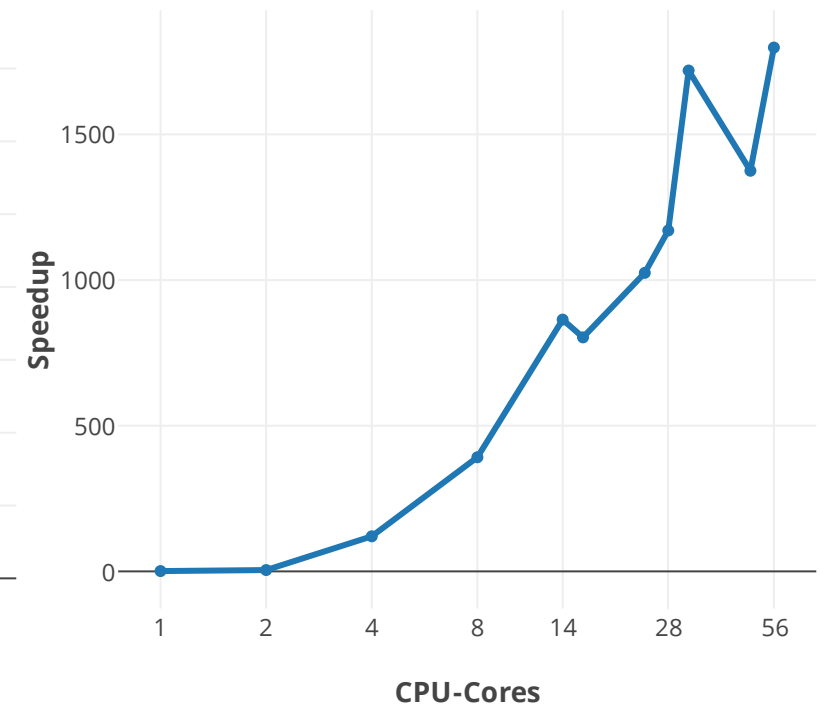
Scaling of C++ Version 1



Scaling of C++ Version 2



Scaling of C# Version 2

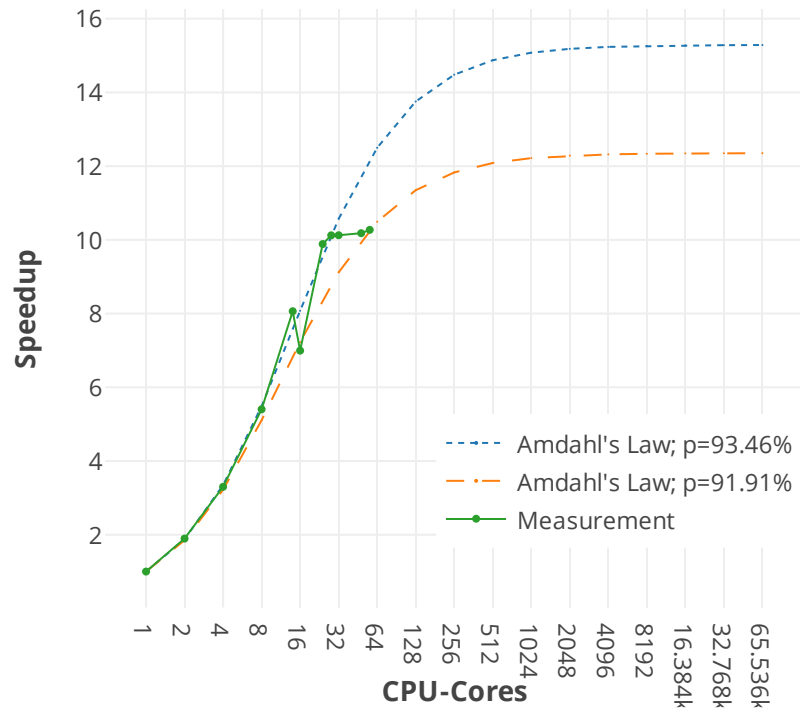




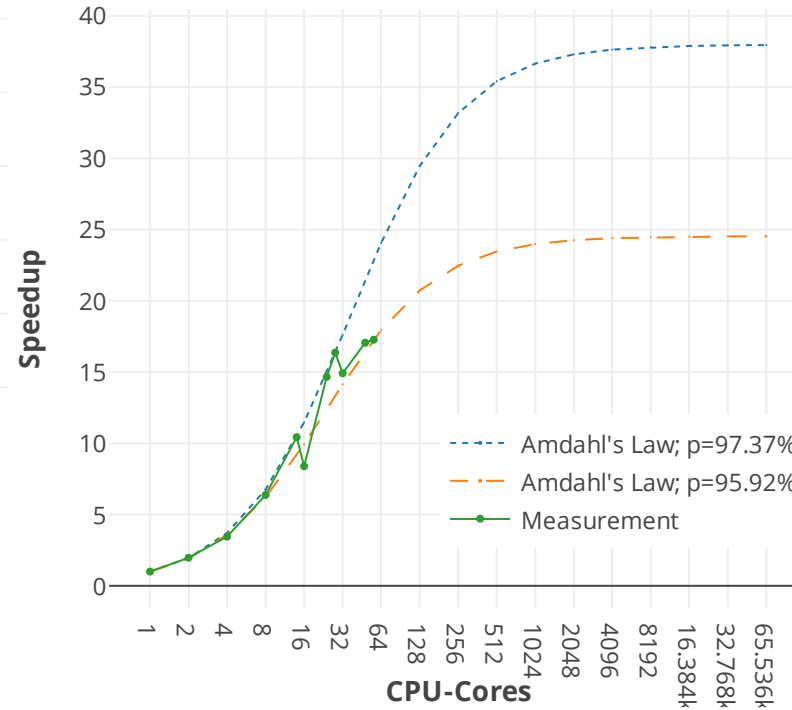
# Scaling with Amdahl's Law of all Versions

## The Himeno Benchmark

### Scaling of C++ Float 32 Version



### Scaling of C++ Float 64 Version



### Scaling of Rust Float 64 Version

